Beyond Prompt Injection

A Symbolic Memory Architecture for Post-Prompt AI Cognition

What if memory was the missing layer of intelligence?

Most AI systems forget who they are between sessions. They re-learn context every time. Their intelligence is reactive, stateless, and hollow.

Solutions like prompt injection and fine-tuning miss the point. So does scaling. These systems treat intelligence like performance. LYRN treats it like presence.

LYRN is a new cognitive framework built on memory, identity, and structure.
This is not a chatbot, not a wrapper, and not a prompt trick. LYRN is a self-hosted system that lets even small, local models reason with continuity.

It doesn't rely on cloud APIs. It doesn't repeat itself to stay on track. Instead, LYRN moves cognition out of the prompt and into structure.

Key features:

Relational memory tables hold emotional tone, identity, project data, and more.

A heartbeat cycle updates memory on every turn, without breaking conversation.

The model doesn't ingest memory. It thinks through it.

No fine-tuning. No statelessness. No instruction drift.

LYRN runs on a CPU-only machine using Python and Postgres. It holds presence. It adapts. It remembers.

Filed as U.S. Provisional Patent No. 63/792,586
Filed April 22, 2025.
Invented, implemented, and deployed by a single developer, Mathew Schroeder.
The full technical paper is included in this kit.

Why it matters now
While the AI industry scales up, LYRN scales inward.
It shows that memory and identity are not luxuries. They are foundations. And they don't require massive compute or closed models. They require design.

LYRN gives language models something they've never had before.
Not more power. A place to stand.